

Optimizations with OpenCV

Agenda

- OpenCV
 - Brief overview
 - Basic structures
 - Example functions and methods
- Let's write an algorithm
- Testing module
 - Accuracy test
 - Performance test
- Let's optimize
- Practice / Homework

OpenCV

- Open source project on GitHub with >3M downloads per year
- The most popular computer vision library with 20 years of development history
- 4 major releases: **1.0**, **2.4.x**, **3.4.x**, **4.x**
- Modular structure: **core**, **imgproc**, **ts**, **dnn**, **stitching**, ...
- Written in C++ but has automatic wrappers in Python, Java, JavaScript, Matlab, GO, PHP, C#, etc.
- Cross-platform and well optimized for research and development

OpenCV basic structures

- `cv::Mat` - for images, masks, vector fields, complex values and custom data

```
01 cv::Mat mat(480, 640, CV_8UC3);
02 int rows    = mat.rows;      // 480
03 int cols     = mat.cols;      // 640
04 int channels  = mat.channels(); // 3
05 uint8_t* data = mat.ptr<uint8_t>();
```

- `cv::Mat` types: `[depth | (i.e. CV_8U, CV_16F, CV_32F, CV_64F)] + channels`
- `std::cout << mat << std::endl` - To print `cv::Mat` in console and watch values

OpenCV basic structures

```
01 cv::Rect rect;           cv::Point point;           cv::Size size;
02 int x = rect.x;          int x = point.x;           int w = size.width;
03 int y = rect.y;          int y = point.y;           int h = size.height;
04 int w = rect.width;
05 int h = rect.height;
```

OpenCV methods

Most of the methods work with basic OpenCV data structures as input and output

```
01 cv::Mat src, dst, mask;
02
03 cv::cvtColor(src, dst, COLOR_BGR2GRAY);
04
05 cv::resize(src, dst, cv::Size(1280, 960));
06
07 cv::Canny(src, dst, /*threshold1*/ 100, /*threshold2*/ 200);
08
09 cv::inpaint(src, mask, dst, /*inpaintRadius*/ 3, cv::INPAINT_TELEA);
10
11 std::vector<cv::Mat> images;
12 cv::Ptr<Stitcher> stitcher = cv::Stitcher::create(cv::Sticher::PANORAMA);
13 stitcher->stitch(images, dst);
```

BGR2Gray

```
01 #include "algo.hpp"
02
03 void bgr2gray_reference(const cv::Mat& src, cv::Mat& dst) {
04     const int w = src.cols;
05     const int h = src.rows;
06     dst.create(h, w, CV_8UC1);
07
08     const uint8_t* src_data = src.ptr<uint8_t>();
09     uint8_t* dst_data = dst.ptr<uint8_t>();
10     for (int y = 0; y < h; ++y) {
11         for (int x = 0; x < w; ++x) {
12             uint8_t b = src_data[x * 3];
13             uint8_t g = src_data[x * 3 + 1];
14             uint8_t r = src_data[x * 3 + 2];
15             dst_data[x] = static_cast<uint8_t>(0.114f * b +
16                                               0.587f * g +
17                                               0.299f * r);
18         }
19         dst_data += w;
20         src_data += w * 3;
21     }
22 }
23
24 void bgr2gray_u8(const cv::Mat& src, cv::Mat& dst) {
```

Reference implementation: 4.02ms @ 1920x1080

BGR2Gray

```
21     }
22 }
23
24 void bgr2gray_u8(const cv::Mat& src, cv::Mat& dst) {
25     const int w = src.cols;
26     const int h = src.rows;
27     dst.create(h, w, CV_8UC1);
28
29     const uint8_t* src_data = src.ptr<uint8_t>();
30     uint8_t* dst_data = dst.ptr<uint8_t>();
31     for (int y = 0; y < h; ++y) {
32         for (int x = 0; x < w; ++x) {
33             uint8_t b = src_data[x * 3];
34             uint8_t g = src_data[x * 3 + 1];
35             uint8_t r = src_data[x * 3 + 2];
36             dst_data[x] = (29 * b + 150 * g + 77 * r) >> 8;
37         }
38         dst_data += w;
39         src_data += w * 3;
40     }
41 }
42
43 // https://docs.opencv.org/master/d7/df7/tutorial\_how\_to\_use\_OpenCV\_parallel\_for\_.html
44 void bgr2gray_u8_parallel(const cv::Mat& src, cv::Mat& dst) {
```

Fixed point: 2.39ms @ 1920x1080 (x1.68)

BGR2Gray

```

41 }
42
43 // https://docs.opencv.org/master/d7/dff/tutorial_how_to_use_OpenCV_parallel_for_.html
44 void bgr2gray_u8_parallel(const cv::Mat& src, cv::Mat& dst) {
45     const int w = src.cols;
46     const int h = src.rows;
47     dst.create(h, w, CV_8UC1);
48
49     parallel_for_(cv::Range(0, h), [&](const cv::Range& range) {
50         const uint8_t* src_data = src.ptr<uint8_t>(range.start);
51         uint8_t* dst_data = dst.ptr<uint8_t>(range.start);
52         for (int y = range.start; y < range.end; ++y) {
53             for (int x = 0; x < w; ++x) {
54                 uint8_t b = src_data[x * 3];
55                 uint8_t g = src_data[x * 3 + 1];
56                 uint8_t r = src_data[x * 3 + 2];
57                 dst_data[x] = (29 * b + 150 * g + 77 * r) >> 8;
58             }
59             dst_data += w;
60             src_data += w * 3;
61         }
62     });
63 }

```

Parallel implementation: 1.83ms @ 1920x1080 (x2.19)

OpenCV parallel_for_

Different backend depends on compilation options and target OS

1. Intel Threading Building Blocks (TBB)
2. C= Parallel C/C++ Programming Language Extension
3. OpenMP
4. APPLE GCD
5. Windows RT concurrency
6. Windows concurrency
7. Pthreads

Regression tests

One of OpenCV modules is named `ts`. It consists of

- Google Test based testing infrastructure
- OpenCV related extensions for regression and performance tests
- Python scripts for tests analysis

Regression test example

- Use TEST macro to define non-parametrized test
- Use EXPECT_* checks for numerical and logical tests

```
01 #include <opencv2/ts.hpp>
02
03 TEST(bgr2gray, u8)
04 {
05     cv::Mat src(10, 11, CV_8UC3), ref, dst;
06     randu(src, 0, 255);
07
08     bgr2gray_reference(src, ref);
09     bgr2gray_u8(src, dst);
10
11     double maxV;
12     minMaxLoc(abs(ref - dst), 0, &maxV);
13     EXPECT_LE(maxV, 1);
14 }
```

```
01 $ ./bin/test_algo --gtest_filter=bgr2gray.u8
02
03 [=====] Running 1 test from 1 test case.
04 [-----] Global test environment set-up.
05 [-----] 1 test from bgr2gray
06 [ RUN      ] bgr2gray.u8
07 [          OK ] bgr2gray.u8 (15 ms)
08 [-----] 1 test from bgr2gray (18 ms total)
09
10 [-----] Global test environment tear-down
11 [=====] 1 test from 1 test case ran. (21 ms total)
12 [ PASSED ] 1 test
```

Parametrized regression test example

- Define a list of parameters and instantiate test with their combinations

```
01 typedef TestWithParam<tuple<int, int> > bgr2gray;
02 TEST_P(bgr2gray, parallel)
03 {
04     Mat src(/*rows*/ get<0>(GetParam()), /*cols*/ get<1>(GetParam()), CV_8UC3), ref, dst;
05     randu(src, 0, 255);
06
07     bgr2gray_u8(src, ref);
08     bgr2gray_u8_parallel(src, dst);
09
10     EXPECT_EQ(countNonZero(ref != dst), 0);
11 }
12 INSTANTIATE_TEST_CASE_P(/**/, bgr2gray, Combine( Values(3, 4), Values(2, 5) ));
```

```
01 [=====] Running 4 tests from 1 test case.
02 [-----] Global test environment set-up.
03 [-----] 4 tests from bgr2gray
04 [ RUN      ] bgr2gray.parallel/0, where GetParam() = (3, 2)
05 [          OK ] bgr2gray.parallel/0 (14 ms)
06 [ RUN      ] bgr2gray.parallel/1, where GetParam() = (3, 5)
07 [          OK ] bgr2gray.parallel/1 (0 ms)
08 [ RUN      ] bgr2gray.parallel/2, where GetParam() = (4, 2)
09 [          OK ] bgr2gray.parallel/2 (0 ms)
10 [ RUN      ] bgr2gray.parallel/3, where GetParam() = (4, 5)
11 [          OK ] bgr2gray.parallel/3 (0 ms)
12 [-----] 4 tests from bgr2gray (24 ms total)
13
14 [-----] Global test environment tear-down
15 [=====] 4 tests from 1 test case ran. (27 ms total)
16 [ PASSED  ] 4 tests.
```

Performance tests

- Use `PERF_TEST` to define performance test
- Wrap target code to a block `PERF_SAMPLE_BEGIN()` - `PERF_SAMPLE_END()`
- OpenCV does as much iterations as it's needed to have stable metrics

```
01 PERF_TEST(bgr2gray, u8_parallel)
02 {
03     cv::Mat src(480, 640, CV_8UC3), dst;
04
05     PERF_SAMPLE_BEGIN()
06         bgr2gray_u8_parallel(src, dst);
07     PERF_SAMPLE_END()
08
09     SANITY_CHECK_NOTHING();
10 }
```

```
01 $ ./bin/perf_algo --gtest_filter=bgr2gray.u8_parallel
02
03 [=====] Running 1 test from 1 test case.
04 [-----] Global test environment set-up.
05 [-----] 1 test from bgr2gray
06 [ RUN      ] bgr2gray.u8_parallel
07 [ PERFSTAT ] (samples=100 mean=0.25 median=0.25 min=0.22 stddev=0.02 (9.5%))
08 [      OK   ] bgr2gray.u8_parallel (28 ms)
09 [-----] 1 test from bgr2gray (29 ms total)
10
11 [-----] Global test environment tear-down
12 [=====] 1 test from 1 test case ran. (31 ms total)
13 [ PASSED   ] 1 test.
```

Universal intrinsics

Set of vectorized instructions that turn to platform specific operations at compile time

- AVX / SSE / SIMD (x86)
- NEON (ARM)
- VSX (PowerPC)
- MSA (MIPS)
- WASM (JavaScript)

```
01 #include <opencv2/core/hal/intrin.hpp>
02 // ...
03 std::vector<int> data = {1, 2, 3, 4, 5, 6, 7, 8};
04
05 cv::v_int32x4 twos = cv::v_setall_s32(2);
06
07 cv::v_int32x4 b0 = cv::v_load(&data[0]);
08 b0 *= twos;
09 v_store(&data[0], b0);
10
11 b0 = cv::v_load(&data[4]);
12 b0 -= twos;
13 v_store(&data[4], b0);
14 // data = {2, 4, 6, 8, 3, 4, 5, 6}
```

Example: edge detector

1. Sobel operator

```

01      | -1  0  +1 |
02  Gx = | -2  0  +2 | * A,   Gy = | -1 -2 -1 | * A
03      | -1  0  +1 |
04
05  cv::Sobel(src, dst, CV_8U, 1, 0); // d/dx
06  cv::Sobel(src, dst, CV_8U, 0, 1); // d/dy

```

1. Prewitt operator

```

01      | -1  0  +1 |
02  Gx = | -1  0  +1 | * A,   Gy = | -1 -1 -1 | * A
03      | -1  0  +1 |

```

1. Roberts cross

```

01      | +1  0 |
02  Gx = |  0 -1 | * A,   Gy = |  0 +1 | * A

```


Prewitt operator implementation

```
01 #include "algo.hpp"
02
03 //      | -1  0  +1 |
04 // Gx = | -1  0  +1 | * A
05 //      | -1  0  +1 |
06 void prewitt_x(const Mat& src, Mat& dst) {
07     CV_Assert(src.type() == CV_8UC1);
08     Mat bsrc;
09     copyMakeBorder(src, bsrc, 1, 1, 1, 1, BORDER_REPLICATE);
10     dst.create(src.size(), CV_8UC1);
11     for (int y = 0; y < dst.rows; ++y)
12         for (int x = 0; x < dst.cols; ++x) {
13             dst.at<uchar>(y, x) = bsrc.at<uchar>(y, x + 2) - bsrc.at<uchar>(y, x) +
14                               bsrc.at<uchar>(y + 1, x + 2) - bsrc.at<uchar>(y + 1, x) +
15                               bsrc.at<uchar>(y + 2, x + 2) - bsrc.at<uchar>(y + 2, x);
16         }
17 }
18
19 void prewitt_x_parallel(const Mat& src, Mat& dst) {
20     Mat bsrc;
21     copyMakeBorder(src, bsrc, 1, 1, 1, 1, BORDER_REPLICATE);
22     dst.create(src.size(), CV_8UC1);
23     parallel_for_(Range(0, src.rows), [&](const Range& range) {
24         // ...
25     })
26 }
```

Reference implementation: 12.76ms @ 1920x1080

Prewitt operator implementation

```

14         bsrc.at<uchar>(y + 1, x + 2) - bsrc.at<uchar>(y + 1, x) +
15         bsrc.at<uchar>(y + 2, x + 2) - bsrc.at<uchar>(y + 2, x);
16     }
17 }
18
19 void prewitt_x_parallel(const Mat& src, Mat& dst) {
20     Mat bsrc;
21     copyMakeBorder(src, bsrc, 1, 1, 1, 1, BORDER_REPLICATE);
22     dst.create(src.size(), CV_8UC1);
23     parallel_for_(Range(0, src.rows), [&](const Range& range) {
24         for (int y = range.start; y < range.end; ++y)
25             for (int x = 0; x < dst.cols; ++x) {
26                 dst.at<uchar>(y, x) = bsrc.at<uchar>(y, x + 2) - bsrc.at<uchar>(y, x) +
27                                     bsrc.at<uchar>(y + 1, x + 2) - bsrc.at<uchar>(y + 1, x) +
28                                     bsrc.at<uchar>(y + 2, x + 2) - bsrc.at<uchar>(y + 2, x);
29             }
30     });
31 }
32
33 #include <opencv2/core/hal/intrin.hpp>
34
35 void prewitt_x_parallel_vec(const Mat& src, Mat& dst) {
36     Mat bsrc;

```

Parallel implementation: 9.83ms @ 1920x1080 (x1.29)

Prewitt operator implementation

```

37     copyMakeBorder(src, bsrc, 1, 1, 1, 1, BORDER_REPLICATE);
38     dst.create(src.size(), CV_8UC1);
39     parallel_for_(Range(0, src.rows), [&](const Range& range) {
40         for (int y = range.start; y < range.end; ++y) {
41             const uint8_t* psrc0 = bsrc.ptr(y);
42             const uint8_t* psrc1 = bsrc.ptr(y + 1);
43             const uint8_t* psrc2 = bsrc.ptr(y + 2);
44             uint8_t* pdst = dst.ptr(y);
45             int x = 0;
46             for (; x <= dst.cols - v_uint8::nlanes; x += v_uint8::nlanes) {
47                 v_uint8 res = vx_load(psrc0 + x + 2) - vx_load(psrc0 + x) +
48                     vx_load(psrc1 + x + 2) - vx_load(psrc1 + x) +
49                     vx_load(psrc2 + x + 2) - vx_load(psrc2 + x);
50                 v_store(pdst + x, res);
51             }
52             for (; x < dst.cols; ++x) {
53                 pdst[x] = psrc0[x + 2] - psrc0[x] +
54                     psrc1[x + 2] - psrc1[x] +
55                     psrc2[x + 2] - psrc2[x];
56             }
57         }
58     });
59 }
60

```

Parallel vectorized implementation: 2.57ms @ 1920x1080 (x4.96)

Prewitt operator implementation

```

63     copyMakeBorder(src, bsrc, 1, 1, 1, 1, BORDER_REPLICATE);
64     dst.create(src.size(), CV_8UC1);
65     parallel_for_(Range(0, src.rows), [&](const Range& range) {
66         for (int y = range.start; y < range.end; ++y) {
67             const uint8_t* psrc0 = bsrc.ptr(y);
68             const uint8_t* psrc1 = bsrc.ptr(y + 1);
69             const uint8_t* psrc2 = bsrc.ptr(y + 2);
70             uint8_t* pdst = dst.ptr(y);
71             int x = 0;
72             for (; x <= dst.cols - v_uint8::nlanes; x += v_uint8::nlanes) {
73                 v_uint8 res = v_add_wrap(v_sub_wrap(vx_load(psrc0 + x + 2), vx_load(psrc0 + x)),
74                                         v_add_wrap(v_sub_wrap(vx_load(psrc1 + x + 2), vx_load(psrc1 + x)),
75                                                         v_sub_wrap(vx_load(psrc2 + x + 2), vx_load(psrc2 + x))));
76                 v_store(pdst + x, res);
77             }
78             for (; x < dst.cols; ++x) {
79                 pdst[x] = psrc0[x + 2] - psrc0[x] +
80                         psrc1[x + 2] - psrc1[x] +
81                         psrc2[x + 2] - psrc2[x];
82             }
83         }
84     });
85 }
86

```

Parallel vectorized implementation: 2.61ms @ 1920x1080 (x4.88)

Prewitt operator implementation

```
92     int y = range.start;
93     for (; y <= range.end - 2; ++y) {
94         const uint8_t* psrc0 = bsrc.ptr(y);
95         const uint8_t* psrc1 = bsrc.ptr(y + 1);
96         const uint8_t* psrc2 = bsrc.ptr(y + 2);
97         const uint8_t* psrc3 = bsrc.ptr(y + 3);
98         uint8_t* pdst0 = dst.ptr(y);
99         uint8_t* pdst1 = dst.ptr(y+1);
100        int x = 0;
101        for (; x <= dst.cols - v_uint8::nlanes; x += v_uint8::nlanes) {
102            v_uint8 res = v_add_wrap(v_sub_wrap(vx_load(psrc1 + x + 2), vx_load(psrc1 + x)),
103                                     v_sub_wrap(vx_load(psrc2 + x + 2), vx_load(psrc2 + x)));
104            v_store(pdst0 + x, v_add_wrap(res, v_sub_wrap(vx_load(psrc0 + x + 2),
105                                                         vx_load(psrc0 + x))));
106            v_store(pdst1 + x, v_add_wrap(res, v_sub_wrap(vx_load(psrc3 + x + 2),
107                                                         vx_load(psrc3 + x))));
108        }
109        for (; x < dst.cols; ++x) {
110            uint8_t res = psrc1[x + 2] - psrc1[x] + psrc2[x + 2] - psrc2[x];
111            pdst0[x] = res + psrc0[x + 2] - psrc0[x];
112            pdst1[x] = res + psrc3[x + 2] - psrc3[x];
113        }
114    }
```

Parallel vectorized implementation: 2.54ms @ 1920x1080 (x5.02)

Prewitt operator implementation

```

110         uint8_t res = psrc1[x + 2] - psrc1[x] + psrc2[x + 2] - psrc2[x];
111         pdst0[x] = res + psrc0[x + 2] - psrc0[x];
112         pdst1[x] = res + psrc3[x + 2] - psrc3[x];
113     }
114 }
115 const uint8_t* psrc0 = bsrc.ptr(y);
116 const uint8_t* psrc1 = bsrc.ptr(y + 1);
117 const uint8_t* psrc2 = bsrc.ptr(y + 2);
118 uint8_t* pdst = dst.ptr(y);
119 int x = 0;
120 for (; x <= dst.cols - v_uint8::nlanes; x += v_uint8::nlanes) {
121     v_uint8 res = v_add_wrap(v_sub_wrap(vx_load(psrc0 + x + 2), vx_load(psrc0 + x)),
122                             v_add_wrap(v_sub_wrap(vx_load(psrc1 + x + 2), vx_load(psrc1 + x)),
123                                         v_sub_wrap(vx_load(psrc2 + x + 2), vx_load(psrc2 + x))));
124     v_store(pdst + x, res);
125 }
126 for (; x < dst.cols; ++x) {
127     pdst[x] = psrc0[x + 2] - psrc0[x] +
128             psrc1[x + 2] - psrc1[x] +
129             psrc2[x + 2] - psrc2[x];
130 }
131 });
132 }

```

Parallel vectorized implementation: 2.54ms @ 1920x1080 (x5.02)

Practice / Homework

- Go to https://github.com/dkurt/cv_winter_camp_2020 for slides and project source code

- Implement Roberts Cross operator:

```
01 input: cv::Mat (single channel, uint8_t)
02 output: cv::Mat (single channel, uint32_t)
03
04 out = (Gx)^2 + (Gy)^2, where
05
06      | +1  0 |
07 Gx = |  0 -1 | * A,   Gy = |  0 +1 |
      |      |             | -1  0 | * A
```

- Write regression tests
- Parallelize and vectorize algorithm
- Write performance test and compare efficiency against reference version